

Homework 4

*Professor: Hank Hoffmann**TA: Will Kong*

1 Due Dates:

To accomodate all the crazy constraints on staffing, the University requirement to submit graduating student grades 10 days before all others, and the problems with SLURM, I am trying something different this time.

The theory problems are due 6/5/2019 at 11:59 PM. This is the latest I can possibly make them due for graduating students and I want everyone to be evaluated on the same criteria.

The coding portion and corresponding writeup of the assignment will be extra credit. For anyone that turns in the coding portion (and the writeup, of course), I will use it to replace your lowest grade. In addition to making this assignment extra credit, I am requiring fewer hash table designs than prior years. Since this portion is extra credit, I will not hold a design review in class. If a sufficient number of people are interested in both the extra credit and a design review, I will hold an extra session for that purpose.

The final writeup will be due 6/14/2019 at 11:59 PM. If you are graduating, you can also do the extra credit, but you will have to turn it in by the same time as the theory HW (sorry, I don't make the rules).

Theory

From the book: 108, 109, 121, 124, 125, 127, 129, 132, 159, 160, 185, 186, 188, 192.

Programming

This assignment is all about high-performance concurrent hash tables. You will build a hash table and extend it in some non-trivial way. Each design will support `ADD()`, `REMOVE()` and `CONTAINS()` methods (as described on page 299 in the text) and will be resizable, based on a *fullness* criteria. That is, you will specify some measure of the occupancy of the table that will trigger a `RESIZE()` operation, which will then double the space and copy the data. There are many such criteria (*e.g.* max total number of items in the table, max total items in any given bucket etc.) and for this assignment you will decide which heuristic to use and provide the reasoning and experimental data (of your own design) to support your decision.

You're permitted to use your lock designs, you may also build on the pthread `MUTEX` provided in the pthreads library. In general, we will use the `STRIPEDHASHSET` paradigm (Figure 13.6 in the text), allocating a lock bank of size $2^{\lceil \log_2 n \rceil}$ (the next larger power of two) for n threads. This allows us to map a hash index to a particular lock with a bitwise AND operation, which is cheap and convenient. Further, in practice it is generally accepted that if the number of locks is approximately equal to the number of threads (with a uniform distribution of accesses to each) you will largely mitigate lock contention. Feel free to demonstrate to yourself that this is true in our case, but it is not necessary for this assignment; you are free to assume it. In addition, we'll make the simplifying assumption that all hash tables have 2^k buckets for some k —after all, we're not at a state school.

Hash Tables We will build each of the following types of hash table, H :

- **Fine Lock Closed-Address (LOCKING):** Use a readers/writer lock to ensure correct access during `ADD()` (including a possible `RESIZE()`), `REMOVE()`, and `CONTAINS()`. This is a standard

hash table where `ADD()` and `REMOVE()` methods use a `WRITELOCK()` to make modifications to the list located at a bucket and `CONTAINS()` uses the corresponding `READLOCK` (see Section 8.3 in the text for a review of the `READWRITELOCK` concept). A `RESIZE()` method merely grabs all `WRITELOCKS` in sequential order (to avoid deadlock with a rival `RESIZE()` attempt) to halt activity during the course of the `RESIZE()` operation.

- **Your Design!** (AWESOME): This is your opportunity to design a really fancy version of an existing idea or try something completely different. In particular, you should extend the design in at least one non-trivial way, demonstrating through experimentation of your own design why it is a good choice and providing your reasoning for correctness. *Please include a section in your report where you concretely discuss the performance tradeoffs you made, your experiments to test those tradeoffs and your reasoning about the correctness of your algorithm.* You are free to pursue whatever design enhancements you like provided they are roughly as sophisticated as the following examples:

- **Optimistic Locking:** This is same as `LOCKING`, except that the `CONTAINS()` method first proceeds without grabbing the corresponding `READLOCK()` and only repeats under the `READLOCK()` if it fails initially. This design requires that any intermediate state during an `ADD()`, `REMOVE()` or `RESIZE()` call must be valid for a concurrent `CONTAINS()` method call.
- **Concurrent Resize:** Design your hash table such that all operations, including `ADD()` and `REMOVE()`, proceed concurrently with `RESIZE()` method calls.
- **Distributed Resize:** When a single thread performs the resize operation, there is the potential that if the rate of all threads making `ADD()` calls exceeds the rate of a single thread copying items, that the maximum speedup is governed by the single thread performing the `RESIZE()` operation. Instead, your design could offload the work of copying items to the other threads. Use your engineering judgment to decide among the myriad ways there are to do this—be sure to include your justification for both performance and correctness in the report.
- **Completely Lock-Free:** Can you build a concurrent hash table using only synchronization primitives like `GETANDSET()` and `FETCHANDINCREMENT` with no locks?
- **Perfect Hashing:** One disadvantage of using a list for each bucket in a Closed-Address hash table is that the traversal of that list is inherently serial and takes time linear in the length of the list in the worst case. In the theoretical analysis of hash tables, there is a technique called **Perfect Hashing** which reduces the access time of a hash table to $O(1)$. Essentially, the technique amounts to another hash table (albeit much smaller) to hold items at each bucket, rather than a list: a hash table of hash tables.

Performance Tests We will test the concurrency of your hash table designs with a simple throughput test, taking packets from a `HASHPACKETGENERATOR` object which are encoded as `ADD`, `REMOVE` or `CONTAINS` requests. Each packet is generated from a source by the Dispatcher, inserted into a queue, and dequeued by a worker. The worker then applies to the corresponding operation to the hash table and calculates the `FINGERPRINT()` as before. We will build three versions of this code:

- **SERIAL:** A single thread serially retrieves packets and makes the appropriate calls into a serial hash table implementation. This version of the code is configurable by the following parameters:
 - `NUMMILLISECONDS` (M) the time in milliseconds that the experiment should run.
 - `FRACTIONADD` (P_+) the percentage of packets that result in an `ADD()` call
 - `FRACTIONREMOVE` (P_-) the fraction of packets that result in a `REMOVE()` call. Note that the fraction of `CONTAINS()` packets is $1 - P_+ - P_-$.
 - `HITRATE` (ρ) the fraction of `CONTAINS()` calls that find data in the hash table.

- `MEAN` (W) the expected amount of work per packet.
- `INITSIZE` (s) the number of items to be preloaded into the table before the Dispatcher and Workers begin working (call `GETADDPACKET()` using the `HASHPACKETGENERATOR` s times).
- **PARALLEL:** As in the previous two assignments, a Dispatcher thread will solely retrieve packets from the `HASHPACKETGENERATOR` and write them into a bank of Lamport queues. Using your best load balancing strategy from the last assignment, the Worker threads will retrieve the packets from the queue bank and then make concurrent calls into the hash table. This version of the code is configurable by the same parameters as `SERIAL` plus:
 - `NUMWORKERS` (n) the number of Worker threads.
 - `TABLETYPE` (H) the type of hash table (*eg.* `LOCKING`, `AWESOME` etc.)
- **PARALLELNOLOAD:** This application is like `PARALLEL`, except that the Workers just drop the packets. The purpose of this application is to determine the maximum throughput of the Dispatcher. It is parameterized by the same parameters as `PARALLEL` minus H .

Experiment

The utilities code for this class contains a `HASHPACKETGENERATOR` you should use for this assignment.

Next, you will perform the following set of experiments across various cross-products of these parameters. Each data point is a measurement taken on a dynamic system (a computer...) and is thus subject to noise. As a result, some care should be taken to extract representative data—we would propose running some reasonable number of experiments for each data point and selecting the median value as the representative. The median of 5 runs should suffice to get a reasonable estimate of the performance on any given configuration - please use your own engineering judgment to decide how many trials you require. Setting the experiment time M to 2000 (*ie.* 2 seconds) should suffice to warm up the caches for all of the following experiments. In each of the following experiments, we describe an experiment that you should perform and analyze. Please provide your hypotheses for the trends you see, supported by additional experiments of your own design if your hypotheses are not sufficiently substantiated by these experiments.

Hash Table Throughput Tests

1. **Dispatcher Rate** Run `PARALLELNOLOAD` with $n = C$ (where C is the number of cores in your test machine) and $W = 1$ to estimate the rate at which the Dispatcher retrieves packets and writes them into the queue bank. How does this rate compare with the same rate in the first firewall programming assignment?
2. **Parallel Overhead** Run `SERIAL` and `PARALLEL` with $n = 1$ and $W = 4000$ on each hash table type, $H \in \{\text{LOCKING}, \text{AWESOME}\}$, with two load configurations:
 - (a) **Mostly Reads** $P_+ = 0.09$, $P_- = 0.01$, $\rho = 0.9$
 - (b) **Heavy Writes** $P_+ = 0.45$, $P_- = 0.05$, $\rho = 0.9$
3. **Speedup** Run `SERIAL` and `PARALLEL` with $n \in \{1, 2, 4, 8, 16, 32\}$, $W = 4000$, all H and under two load configurations:
 - (a) **Mostly Reads** $P_+ = 0.09$, $P_- = 0.01$, $\rho \in \{0.5, 0.75, 0.9, 0.99\}$
 - (b) **Heavy Writes** $P_+ = 0.45$, $P_- = 0.05$, $\rho \in \{0.5, 0.75, 0.9, 0.99\}$

Make plots (one for each load configuration) with speedup (relative to `SERIAL`) on the Y -axis and n on the X -axis. Analyze this data and present your most insightful observations in the report.

Writeup

Please submit a typeset (learn \LaTeX if you haven't already!) report summarizing your results from the experiments and the conclusions you draw from them. Your report should include the experiments as specified above, any additional experimental data you used to draw conclusions and a self-contained report. That is, somebody should be able to read the report alone and understand what code you developed, what experiments you ran and how the data supports the conclusions you draw. In addition, your report should provide the reasoning (though, mercifully, not a formal proof) as to why your designs are correct. In particular, for each method in each design you should specify whether the method is wait-free, lock-free, deadlock-free or starvation-free and the reasoning for it. Finally, submit your development directory, containing all of the working code.